

NLPDisasterTweets

November 8, 2024

1 Introduction to Deep Learning - Week 4

1.1 NLP Disaster Tweets Kaggle Mini Project

Github Link <https://github.com/conditas/NLPDisasterTweet>

This project uses Natural Language Processing and a Recurrent Neural Network Model (RNN) for the Kaggle Competition “Natural Language Processing with Disaster Tweets.” It involves being able to predict and classify whether twitter texts are about an actual disaster or not. This can be helpful for spreading real-time information during emergency situations.

1.1.1 Import Libraries

The following libraries will be used in my project.

```
[212]: %matplotlib inline
import warnings
warnings.simplefilter("ignore", FutureWarning)

import pandas as pd
import numpy as np
import scipy as sp
import scipy.stats as stats
import seaborn as sns
import matplotlib.pyplot as plt

import gc
import os

import sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score, \
    accuracy_score, ConfusionMatrixDisplay, confusion_matrix, f1_score
from tensorflow.keras.preprocessing.text import Tokenizer

import tensorflow as tf
from tensorflow import keras
from keras import layers, optimizers
```

```

import keras_core as keras
import keras_nlp
from tensorflow.keras.models import Sequential
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.layers import Embedding, LSTM, Bidirectional,
    Dense, GlobalMaxPool1D, BatchNormalization, Dropout
from tensorflow.keras.preprocessing.sequence import pad_sequences

os.environ['KERAS_BACKEND'] = 'tensorflow'

```

1.1.2 Load the Data

The data comes from the “Natural Language Processing with Disaster Tweets” Kaggle competition. It consists of a two files a train.csv and a test.csv. The train file contains 5 columns and 7613 rows. The test file will be used for the kaggle competition submission and has 4 columns with 3263 rows. The table below summarizes the variables and below the table I print out the data types for each column.

Variable	Description
id	unique identifier
keyword	keyword from tweet
location	where tweet was sent from
text	text of twitter tweet
target	classifier ‘1’ for disaster and ‘0’ for non-disaster

Data Citation

Addison Howard, devrishi, Phil Culliton, and Yufeng Guo. Natural Language Processing with Disaster Tweets. <https://kaggle.com/competitions/nlp-getting-started/data>, 2019. Kaggle.

```

[238]: df_train = pd.read_csv("train.csv")
df_test = pd.read_csv("test.csv")

print('Training:\n')
print('Shape = {}'.format(df_train.shape))
print(df_train.dtypes)

print('\nTest:\n')
print('Shape = {}'.format(df_test.shape))
print(df_test.dtypes)

print("Example of Data:\n", df_train.head(5))

```

Training:

```

Shape = (7613, 5)
id          int64
keyword     object

```

```
location    object
text        object
target      int64
dtype: object
```

Test:

```
Shape = (3263, 4)
id      int64
keyword object
location object
text    object
dtype: object
```

Example of Data:

	id	keyword	location	text \
0	1	NaN	NaN	Our Deeds are the Reason of this #earthquake M...
1	4	NaN	NaN	Forest fire near La Ronge Sask. Canada
2	5	NaN	NaN	All residents asked to 'shelter in place' are ...
3	6	NaN	NaN	13,000 people receive #wildfires evacuation or...
4	7	NaN	NaN	Just got sent this photo from Ruby #Alaska as ...

	target
0	1
1	1
2	1
3	1
4	1

1.1.3 Exploratory Data Analysis

In my exploratory data analysis I compare the number of disaster tweets vs non-disaster tweets and calculate an imbalance ratio. This ratio was 1.35 which should be fine for modeling. I check for null values across the columns. There were not any null or empty values in the text column. The distribution of locations shows most of the tweets came from the US. There were some discrepancies on the naming (USA vs United States) and the location column had a significant number of missing values (33%) so I will not be using this information. I compare text lengths of disaster vs non-disaster tweets. In general, the disaster tweets tend to be longer than non-disaster.

Finally I split the training data set 80/20 in order to train and validate models.

```
[239]: text_lengths_1 = []
text_lengths_0 = []

####
#### This code plots text length distribution for disaster and non-disaster
      ↪ tweets
####
for index,row in df_train.iterrows():
```

```

if row['target']==1:
    text_lengths_1.append(len(row['text']))

if row['target']==0:
    text_lengths_0.append(len(row['text']))

sns.histplot(text_lengths_0, label='Not Disaster')
sns.histplot(text_lengths_1, label='Disaster')
plt.title("Train Data - Text Lengths Distribution")
plt.xlabel('Text Length')
plt.legend(loc='upper left')
plt.show()

print("Examples of Disaster Tweets")
print(df_train[df_train['target']==1]['text'][0:10])

print("\nExamples of Non-Disaster Tweets")
print(df_train[df_train['target']==0]['text'][0:10])

####
## Clean data- looking for null values for text column
####
print("train dimensions: ", df_train.shape)
print("test dimensions: ", df_test.shape)

#Checking for null values and duplicates
print("\nTrain Dataset")
print("\nNull values in text?: ",df_train['text'].isnull().values.sum())
print("\nNull values in keyword?: ",df_train['keyword'].isnull().values.sum())
print("\nNull values in location?: ",df_train['location'].isnull().values.sum())
print("\n\nTest Dataset")
print("\nNull values in text?: ",df_test['text'].isnull().values.sum())
print("\nNull values in keyword?: ",df_test['keyword'].isnull().values.sum())
print("\nNull values in location?: ",df_test['location'].isnull().values.sum())

#####
# counts plot by target
####
print()
plt.figure(figsize = (5,5))
sns.histplot(df_train['target'],bins=range(3), ec='k')
plt.xticks([0,1])
plt.title("Target Count Real vs Not Real")
plt.show()

####

```

```

#calculate imbalance ratio
####
num_zeros = (y_train== 0).sum()
num_ones = (y_train == 1).sum()
imbalance = num_zeros/num_ones
print("Imbalance: ", imbalance)

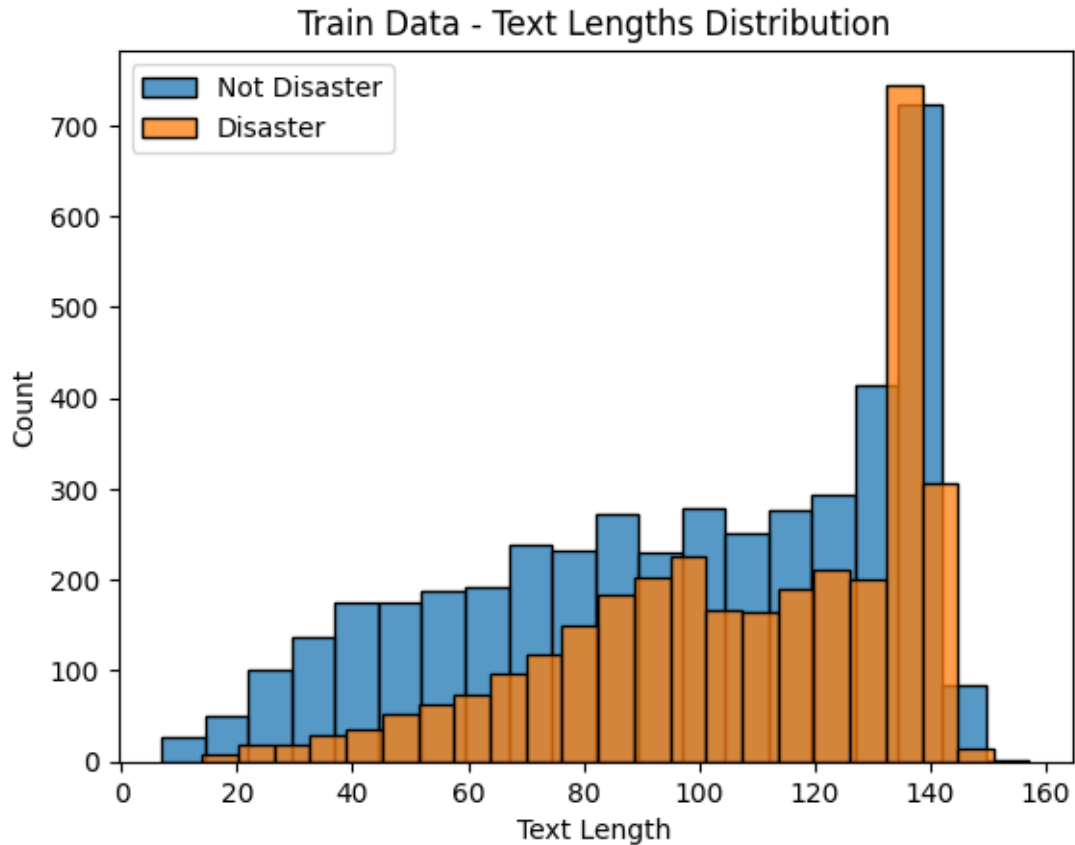
#####
# counts plot by location
#####
print("\nWhere are tweets from?")
print(df_train['location'].value_counts()[:10])
plt.figure(figsize = (5,5))
sns.barplot(y=df_train['location'].value_counts()[:10],
            ↪index,x=df_train['location'].value_counts()[:10], color='maroon')
plt.title("Tweet Locations")
plt.show()

###
## Split into train file into train and validation datasets
###

target = df_train["target"]
x = df_train.drop(["id", "location", "target"], axis=1)
X_train, X_test, y_train, y_test = train_test_split(df_train["text"], target, ↪
            ↪test_size=0.2, random_state=1234)

print("Train Shape X/Y: ")
print(X_train.shape)
print(y_train.shape)
print("Validation Shape X/Y: ")
print(X_test.shape)
print(y_test.shape)

```



Examples of Disaster Tweets

```

0   Our Deeds are the Reason of this #earthquake M...
1       Forest fire near La Ronge Sask. Canada
2   All residents asked to 'shelter in place' are ...
3   13,000 people receive #wildfires evacuation or...
4   Just got sent this photo from Ruby #Alaska as ...
5   #RockyFire Update => California Hwy. 20 closed...
6   #flood #disaster Heavy rain causes flash flood...
7   I'm on top of the hill and I can see a fire in...
8   There's an emergency evacuation happening now ...
9   I'm afraid that the tornado is coming to our a...
Name: text, dtype: object

```

Examples of Non-Disaster Tweets

```

15       What's up man?
16       I love fruits
17       Summer is lovely
18       My car is so fast
19   What a goooooooooaaaaaal!!!!!!
20       this is ridiculous...

```

```
21           London is cool ;)
22           Love skiing
23           What a wonderful day!
24           L000000L
Name: text, dtype: object
train dimensions: (7613, 5)
test dimensions: (3263, 4)
```

Train Dataset

```
Null values in text?: 0
```

```
Null values in keyword?: 61
```


```
Null values in location?: 2533
```

Test Dataset

```
Null values in text?: 0
```

```
Null values in keyword?: 26
```

```
Null values in location?: 1105
```



Imbalance: 1.3522595596755504

Where are tweets from?

location

USA 104

New York 71

United States 50

London 45

Canada 29

Nigeria 28


UK 27

Los Angeles, CA 26

India 24

Mumbai 22

Name: count, dtype: int64



Train Shape X/Y:

(6090,)

(6090,)

Validation Shape X/Y:

(1523,)

(1523,)

1.1.4 Models

The first step in creating the model was to process the text. I used the Keras Tokenizer API which converts all characters to lowercase, it removes unnecessary characters such as symbols, punctuation, and numbers. It converts the text into a list of words in a string. I create a word embedding using GloVe (Global Vectors for Word Representation) for the model. This converts the sequence of words into a vectorized form. This vector format captures semantic relationships between words.

Since sequence in text is important, I use an Recurrent Neural Network model with LSTM (Long Short Term Memory). I use the embedding from GloVe in the input layer to transform the input sequences into a vector of size 100. I then use a bi-directional LSTM layer followed by a dense layer with an relu activation function. I use a drop out layer for regulation and then a final dense layer with a sigmoid activation function for the output label classification.

```
[287]: # Text Preprocessing
# This code performs tokenization of the tweet text
tokenizer = Tokenizer() # # create the tokenizer
tokenizer.fit_on_texts(X_train)
tokenizer.fit_on_texts(X_test)

X_train_seq = tokenizer.texts_to_sequences(X_train)
X_test_seq = tokenizer.texts_to_sequences(X_test)
x_sub_test = tokenizer.texts_to_sequences(df_test["text"])
# Pad sequences to a fixed length
X_train_seq = pad_sequences(X_train_seq, truncating='post', padding='post')
X_test_seq = pad_sequences(X_test_seq, truncating='post', padding='post')
x_sub_test = pad_sequences(x_sub_test, truncating='post', padding='post')
X_train_seq.shape
```

[287]: (6090, 33)

```
[288]: ### Embedding Function Using Glove 6B 100 text file

def embedding_for_vocab(filepath, word_index,
                        embedding_dim):
    vocab_size = len(word_index) + 1

    embedding_matrix_vocab = np.zeros((vocab_size,
                                       embedding_dim))

    with open(filepath, encoding="utf8") as f:
        for line in f:
            word, *vector = line.split()
            if word in word_index:
                idx = word_index[word]
                embedding_matrix_vocab[idx] = np.array(
                    vector, dtype=np.float32)[:embedding_dim]

    return embedding_matrix_vocab

embedding_matrix_vocab = embedding_for_vocab(
    'glove.6B.100d.txt', tokenizer.word_index,
    100)
```

```
[289]: #####
# This code builds the RNN model
#####
opt = optimizers.Adam(learning_rate=0.01, beta_1=0.9)
```

```

model = Sequential()
model.add(Embedding(input_dim=embedding_matrix_vocab.shape[0],
                    output_dim=embedding_matrix_vocab.shape[1],
                    weights=[embedding_matrix_vocab],
                    input_length=30,
                    trainable=False))
model.add(Bidirectional(LSTM(64, recurrent_dropout=0.2)))
model.add(Dropout(0.2))
model.add(Dense(64, activation = "relu"))
model.add(Dropout(0.2))
model.add(Dense(1, activation='sigmoid'))

model.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])

```

```

/Library/Frameworks/Python.framework/Versions/3.11/lib/python3.11/site-
packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument
`input_length` is deprecated. Just remove it.
  warnings.warn(

```

1.1.5 Results and Analysis

I run an initial model and plot the accuracy of the training data compared to the validation data. The accuracy of the training set is pretty decent- 0.98 at the 10th epoch. However the validation accuracy is only 0.80 which suggests there is overfitting in the model. The F1 score was 0.77. I perform some hyperparameter tuning to get this a little higher.

```

[323]: #####
# Train the model
#####
lstm_model = model.fit(X_train_seq, y_train, epochs=10, batch_size=32,
    ↪ validation_data = (X_test_seq, y_test))

plt.title('Accuracy')
plt.plot(lstm_model.history['accuracy'], label='train')
plt.plot(lstm_model.history['val_accuracy'], label='test')
plt.legend()
plt.show();

#predictions and f1 score
y_pred = np.round(model.predict(X_test_seq))

print("F1 Score: ", f1_score(y_test, y_pred))


```

```

Epoch 1/10
191/191          3s 15ms/step -
accuracy: 0.9845 - loss: 0.0296 - val_accuracy: 0.8063 - val_loss: 2.2264
Epoch 2/10

```

191/191 3s 15ms/step -
accuracy: 0.9838 - loss: 0.0302 - val_accuracy: 0.8070 - val_loss: 2.1891
Epoch 3/10
191/191 3s 15ms/step -
accuracy: 0.9849 - loss: 0.0302 - val_accuracy: 0.8024 - val_loss: 2.2214
Epoch 4/10
191/191 3s 15ms/step -
accuracy: 0.9817 - loss: 0.0310 - val_accuracy: 0.8050 - val_loss: 2.2162
Epoch 5/10
191/191 3s 15ms/step -
accuracy: 0.9856 - loss: 0.0390 - val_accuracy: 0.8043 - val_loss: 2.1365
Epoch 6/10
191/191 3s 15ms/step -
accuracy: 0.9825 - loss: 0.0315 - val_accuracy: 0.8089 - val_loss: 2.1296
Epoch 7/10
191/191 3s 15ms/step -
accuracy: 0.9837 - loss: 0.0323 - val_accuracy: 0.8083 - val_loss: 2.1701
Epoch 8/10
191/191 3s 15ms/step -
accuracy: 0.9846 - loss: 0.0291 - val_accuracy: 0.8050 - val_loss: 2.2688
Epoch 9/10
191/191 3s 15ms/step -
accuracy: 0.9825 - loss: 0.0339 - val_accuracy: 0.8030 - val_loss: 2.2499
Epoch 10/10
191/191 3s 15ms/step -
accuracy: 0.9839 - loss: 0.0314 - val_accuracy: 0.8017 - val_loss: 2.1582



48/48 0s 4ms/step
 F1 Score: 0.770516717325228

Hyperparameter Tuning

I tried different approaches to tuning the hyperparameters. I tried lowering the learning rate, changing the optimizer to RMS Prop, increasing batch size and epoch number. Finally I did I combination of increased epochs and decreased learning rate. The results are summarized in a table below.

```
[324]: #####
# Hyperparameter Tuning -
#####

# Decrease Learning Rate by half
opt = optimizers.Adam(learning_rate=0.005, beta_1=0.9)
model.compile(optimizer=opt,loss='binary_crossentropy',metrics=['accuracy'])

lstm_model_2 = model.fit(X_train_seq, y_train, epochs=10, batch_size=32,
    ↪validation_data = (X_test_seq, y_test))
```

```

plt.title('Accuracy Decreased Learning Rate')
plt.plot(lstm_model_2.history['accuracy'], label='train')
plt.plot(lstm_model_2.history['val_accuracy'], label='test')
plt.legend()
plt.show();

#predictions and f1 score
y_pred =np.round(model.predict(X_test_seq))
print("F1 Score Decreased Learning Rate: ",f1_score(y_test, y_pred))


# Change Optimizer from ADAM to RMS Prop

#opt = optimizers.RMSPROP(learning_rate=0.01)
model.
    ↪compile(optimizer='rmsprop',loss='binary_crossentropy',metrics=['accuracy'])

lstm_model_3 = model.fit(X_train_seq, y_train, epochs=10, batch_size=32,
    ↪validation_data = (X_test_seq, y_test))

plt.title('Accuracy- RMS PROP ')
plt.plot(lstm_model_3.history['accuracy'], label='train')
plt.plot(lstm_model_3.history['val_accuracy'], label='test')
plt.legend()
plt.show();

#predictions and f1 score
y_pred =np.round(model.predict(X_test_seq))
print("F1 Score RMSProp Optimizer: ",f1_score(y_test, y_pred))


# Change Batch Size from 32 to 64

opt = optimizers.Adam(learning_rate=0.01, beta_1=0.9)
model.
    ↪compile(optimizer='rmsprop',loss='binary_crossentropy',metrics=['accuracy'])

lstm_model_4 = model.fit(X_train_seq, y_train, epochs=10, batch_size=64,
    ↪validation_data = (X_test_seq, y_test))

plt.title('Accuracy- Increased Batch Size')
plt.plot(lstm_model_4.history['accuracy'], label='train')
plt.plot(lstm_model_4.history['val_accuracy'], label='test')
plt.legend()
plt.show();

```

```

#predictions and f1 score
y_pred =np.round(model.predict(X_test_seq))
print("F1 Score Increased Batch Size: ",f1_score(y_test, y_pred))

# Increase Epoch
opt = optimizers.Adam(learning_rate=0.01, beta_1=0.9)
model.
    ↪compile(optimizer='rmsprop',loss='binary_crossentropy',metrics=['accuracy'])

lstm_model_5 = model.fit(X_train_seq, y_train, epochs=20, batch_size=32,
    ↪validation_data = (X_test_seq, y_test))

plt.title('Accuracy- Increased Epoch Size')
plt.plot(lstm_model_5.history['accuracy'], label='train')
plt.plot(lstm_model_5.history['val_accuracy'], label='test')
plt.legend()
plt.show();

#predictions and f1 score
y_pred =np.round(model.predict(X_test_seq))
print("F1 Score Increased Epoch Number: ",f1_score(y_test, y_pred))

```

```

Epoch 1/10
191/191          4s 12ms/step -
accuracy: 0.9840 - loss: 0.0434 - val_accuracy: 0.7932 - val_loss: 1.8165
Epoch 2/10
191/191          3s 14ms/step -
accuracy: 0.9783 - loss: 0.0570 - val_accuracy: 0.8043 - val_loss: 1.8758
Epoch 3/10
191/191          3s 15ms/step -
accuracy: 0.9795 - loss: 0.0539 - val_accuracy: 0.8011 - val_loss: 1.7285
Epoch 4/10
191/191          3s 15ms/step -
accuracy: 0.9736 - loss: 0.0723 - val_accuracy: 0.8063 - val_loss: 1.4450
Epoch 5/10
191/191          3s 15ms/step -
accuracy: 0.9790 - loss: 0.0513 - val_accuracy: 0.7958 - val_loss: 1.3402
Epoch 6/10
191/191          3s 15ms/step -
accuracy: 0.9769 - loss: 0.0533 - val_accuracy: 0.8037 - val_loss: 1.3474
Epoch 7/10
191/191          3s 15ms/step -
accuracy: 0.9741 - loss: 0.0596 - val_accuracy: 0.8011 - val_loss: 1.3721
Epoch 8/10
191/191          3s 15ms/step -
accuracy: 0.9800 - loss: 0.0502 - val_accuracy: 0.7971 - val_loss: 1.1951

```

Epoch 9/10


191/191 3s 15ms/step -

accuracy: 0.9786 - loss: 0.0459 - val_accuracy: 0.7912 - val_loss: 1.5901

Epoch 10/10

191/191 3s 15ms/step -

accuracy: 0.9765 - loss: 0.0531 - val_accuracy: 0.7945 - val_loss: 1.5667



48/48

0s 5ms/step

F1 Score Decreased Learning Rate: 0.7658937920718025

Epoch 1/10

191/191 3s 12ms/step -

accuracy: 0.9774 - loss: 0.0494 - val_accuracy: 0.8004 - val_loss: 1.6714

Epoch 2/10

191/191 3s 14ms/step -

accuracy: 0.9848 - loss: 0.0350 - val_accuracy: 0.8011 - val_loss: 1.7988

Epoch 3/10

191/191 3s 15ms/step -

accuracy: 0.9829 - loss: 0.0389 - val_accuracy: 0.7997 - val_loss: 1.9420

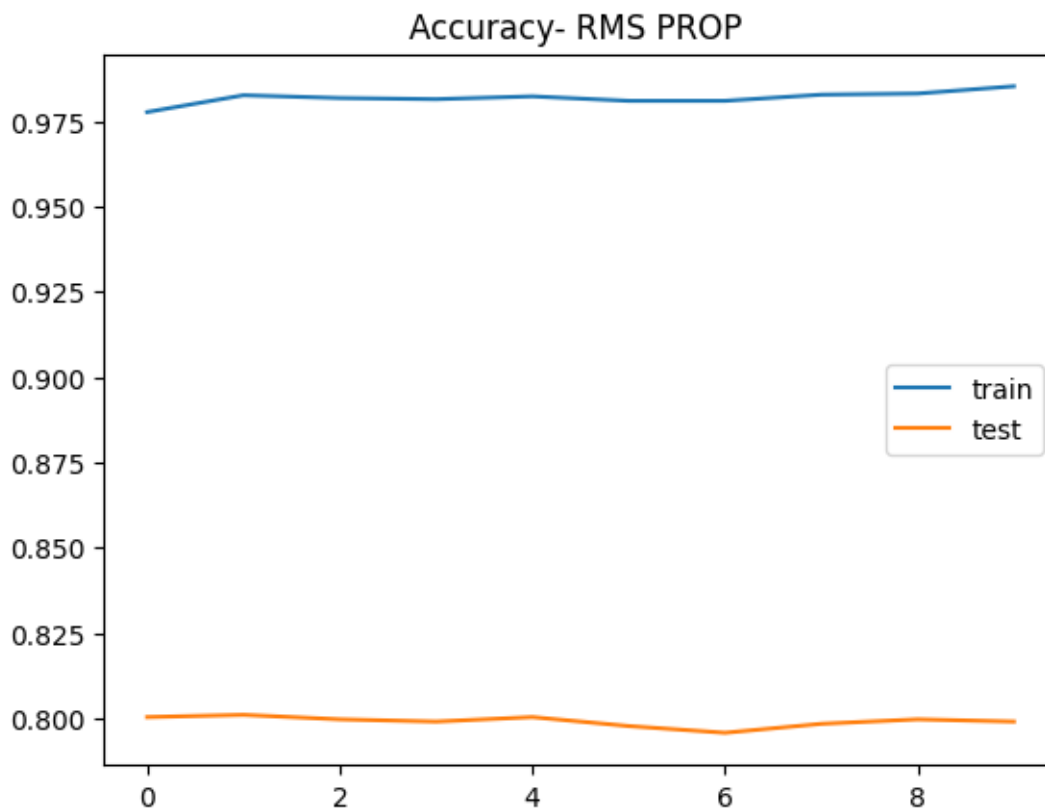
Epoch 4/10

191/191 3s 15ms/step -

accuracy: 0.9831 - loss: 0.0393 - val_accuracy: 0.7991 - val_loss: 1.9633

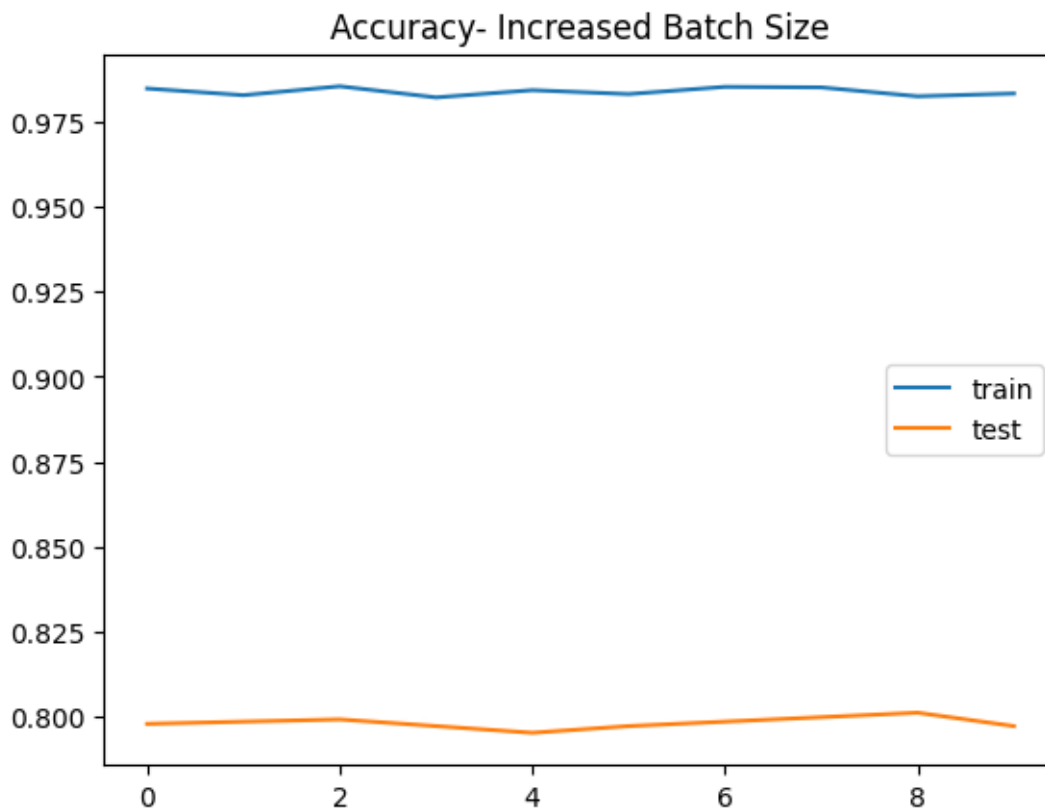
Epoch 5/10

191/191 3s 15ms/step -
accuracy: 0.9840 - loss: 0.0378 - val_accuracy: 0.8004 - val_loss: 1.9944
Epoch 6/10
191/191 3s 15ms/step -
accuracy: 0.9827 - loss: 0.0383 - val_accuracy: 0.7978 - val_loss: 2.0073
Epoch 7/10
191/191 3s 15ms/step -
accuracy: 0.9824 - loss: 0.0375 - val_accuracy: 0.7958 - val_loss: 1.9941
Epoch 8/10
191/191 3s 14ms/step -
accuracy: 0.9833 - loss: 0.0363 - val_accuracy: 0.7984 - val_loss: 2.0600
Epoch 9/10
191/191 3s 15ms/step -
accuracy: 0.9822 - loss: 0.0383 - val_accuracy: 0.7997 - val_loss: 2.0870
Epoch 10/10
191/191 3s 15ms/step -
accuracy: 0.9834 - loss: 0.0355 - val_accuracy: 0.7991 - val_loss: 2.1883



48/48 0s 5ms/step
F1 Score RMSProp Optimizer: 0.7649769585253456
Epoch 1/10
96/96 3s 16ms/step -

accuracy: 0.9856 - loss: 0.0307 - val_accuracy: 0.7978 - val_loss: 2.2161
Epoch 2/10
96/96 2s 24ms/step -
accuracy: 0.9834 - loss: 0.0395 - val_accuracy: 0.7984 - val_loss: 2.1742
Epoch 3/10
96/96 2s 24ms/step -
accuracy: 0.9871 - loss: 0.0316 - val_accuracy: 0.7991 - val_loss: 2.1967
Epoch 4/10
96/96 2s 24ms/step -
accuracy: 0.9822 - loss: 0.0352 - val_accuracy: 0.7971 - val_loss: 2.2207
Epoch 5/10
96/96 2s 24ms/step -
accuracy: 0.9852 - loss: 0.0344 - val_accuracy: 0.7951 - val_loss: 2.1893
Epoch 6/10
96/96 2s 25ms/step -
accuracy: 0.9843 - loss: 0.0319 - val_accuracy: 0.7971 - val_loss: 2.2188
Epoch 7/10
96/96 2s 25ms/step -
accuracy: 0.9865 - loss: 0.0368 - val_accuracy: 0.7984 - val_loss: 2.2432
Epoch 8/10
96/96 2s 25ms/step -
accuracy: 0.9861 - loss: 0.0307 - val_accuracy: 0.7997 - val_loss: 2.2656
Epoch 9/10
96/96 2s 25ms/step -
accuracy: 0.9833 - loss: 0.0356 - val_accuracy: 0.8011 - val_loss: 2.2200
Epoch 10/10
96/96 2s 25ms/step -
accuracy: 0.9821 - loss: 0.0375 - val_accuracy: 0.7971 - val_loss: 2.2704

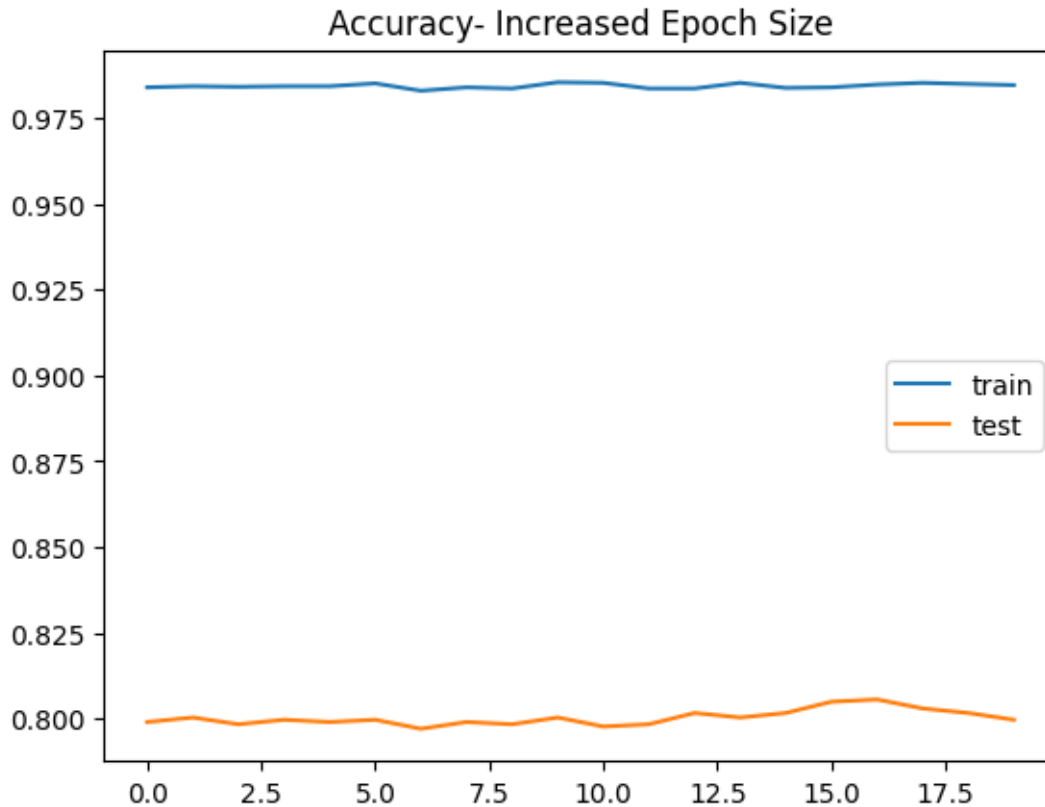


```

48/48          0s 5ms/step
F1 Score Increased Batch Size:  0.7621247113163973
Epoch 1/20
191/191        5s 12ms/step -
accuracy: 0.9835 - loss: 0.0317 - val_accuracy: 0.7991 - val_loss: 2.3515
Epoch 2/20
191/191        3s 14ms/step -
accuracy: 0.9826 - loss: 0.0324 - val_accuracy: 0.8004 - val_loss: 2.4363
Epoch 3/20
191/191        3s 15ms/step -
accuracy: 0.9843 - loss: 0.0353 - val_accuracy: 0.7984 - val_loss: 2.4105
Epoch 4/20
191/191        3s 15ms/step -
accuracy: 0.9859 - loss: 0.0294 - val_accuracy: 0.7997 - val_loss: 2.5282
Epoch 5/20
191/191        3s 15ms/step -
accuracy: 0.9840 - loss: 0.0406 - val_accuracy: 0.7991 - val_loss: 2.4813
Epoch 6/20
191/191        3s 15ms/step -
accuracy: 0.9853 - loss: 0.0367 - val_accuracy: 0.7997 - val_loss: 2.4904
Epoch 7/20

```

191/191 3s 15ms/step -
 accuracy: 0.9820 - loss: 0.0450 - val_accuracy: 0.7971 - val_loss: 2.5011
 Epoch 8/20
 191/191 3s 15ms/step -
 accuracy: 0.9832 - loss: 0.0371 - val_accuracy: 0.7991 - val_loss: 2.5134
 Epoch 9/20
 191/191 3s 15ms/step -
 accuracy: 0.9840 - loss: 0.0325 - val_accuracy: 0.7984 - val_loss: 2.6135
 Epoch 10/20
 191/191 3s 15ms/step -
 accuracy: 0.9856 - loss: 0.0277 - val_accuracy: 0.8004 - val_loss: 2.7104
 Epoch 11/20
 191/191 3s 15ms/step -
 accuracy: 0.9855 - loss: 0.0336 - val_accuracy: 0.7978 - val_loss: 2.8001
 Epoch 12/20
 191/191 3s 15ms/step -
 accuracy: 0.9818 - loss: 0.0352 - val_accuracy: 0.7984 - val_loss: 2.7301
 Epoch 13/20
 191/191 3s 15ms/step -
 accuracy: 0.9826 - loss: 0.0361 - val_accuracy: 0.8017 - val_loss: 2.6874
 Epoch 14/20
 191/191 3s 15ms/step -
 accuracy: 0.9851 - loss: 0.0315 - val_accuracy: 0.8004 - val_loss: 2.6988
 Epoch 15/20
 191/191 3s 15ms/step -
 accuracy: 0.9847 - loss: 0.0276 - val_accuracy: 0.8017 - val_loss: 2.7968
 Epoch 16/20
 191/191 3s 15ms/step -
 accuracy: 0.9805 - loss: 0.0377 - val_accuracy: 0.8050 - val_loss: 2.7952
 Epoch 17/20
 191/191 3s 15ms/step -
 accuracy: 0.9873 - loss: 0.0314 - val_accuracy: 0.8056 - val_loss: 2.7331
 Epoch 18/20
 191/191 3s 15ms/step -
 accuracy: 0.9854 - loss: 0.0347 - val_accuracy: 0.8030 - val_loss: 2.8794
 Epoch 19/20
 191/191 3s 15ms/step -
 accuracy: 0.9844 - loss: 0.0358 - val_accuracy: 0.8017 - val_loss: 2.9339
 Epoch 20/20
 191/191 3s 15ms/step -
 accuracy: 0.9864 - loss: 0.0296 - val_accuracy: 0.7997 - val_loss: 3.0109



48/48 0s 5ms/step
 F1 Score Increased Epoch Number: 0.765564950038432

```
[321]: #Decrease Learning Increase Epoch
opt = optimizers.Adam(learning_rate=0.001, beta_1=0.9)
model.compile(optimizer=opt,loss='binary_crossentropy',metrics=['accuracy'])

lstm_model_6 = model.fit(X_train_seq, y_train, epochs=50, batch_size=32,
    ↪ validation_data = (X_test_seq, y_test))

plt.title('Accuracy- Combined Increased Epoch Size/Decrease Learning Rate')
plt.plot(lstm_model_6.history['accuracy'], label='train')
plt.plot(lstm_model_6.history['val_accuracy'], label='test')
plt.legend()
plt.show();

#predictions and f1 score
y_pred =np.round(model.predict(X_test_seq))
print("F1 Score Combined Increased Epoch Size/Decrease Learning Rate:
    ↪",f1_score(y_test, y_pred))
```


Epoch 1/50

191/191 4s 12ms/step -
 accuracy: 0.9824 - loss: 0.0326 - val_accuracy: 0.8024 - val_loss: 2.3037
 Epoch 2/50
 191/191 3s 14ms/step -
 accuracy: 0.9833 - loss: 0.0333 - val_accuracy: 0.7991 - val_loss: 2.1897
 Epoch 3/50
 191/191 3s 15ms/step -
 accuracy: 0.9847 - loss: 0.0326 - val_accuracy: 0.8017 - val_loss: 2.1412
 Epoch 4/50
 191/191 3s 15ms/step -
 accuracy: 0.9842 - loss: 0.0346 - val_accuracy: 0.8011 - val_loss: 2.1349
 Epoch 5/50
 191/191 3s 15ms/step -
 accuracy: 0.9851 - loss: 0.0311 - val_accuracy: 0.8043 - val_loss: 2.1277
 Epoch 6/50
 191/191 3s 15ms/step -
 accuracy: 0.9823 - loss: 0.0334 - val_accuracy: 0.8004 - val_loss: 2.1378
 Epoch 7/50
 191/191 3s 15ms/step -
 accuracy: 0.9860 - loss: 0.0326 - val_accuracy: 0.8024 - val_loss: 2.1487
 Epoch 8/50
 191/191 3s 15ms/step -
 accuracy: 0.9849 - loss: 0.0362 - val_accuracy: 0.8011 - val_loss: 2.1865
 Epoch 9/50
 191/191 3s 15ms/step -
 accuracy: 0.9841 - loss: 0.0360 - val_accuracy: 0.7984 - val_loss: 2.1513
 Epoch 10/50
 191/191 3s 15ms/step -
 accuracy: 0.9850 - loss: 0.0328 - val_accuracy: 0.7978 - val_loss: 2.2755
 Epoch 11/50
 191/191 3s 15ms/step -
 accuracy: 0.9844 - loss: 0.0332 - val_accuracy: 0.8011 - val_loss: 2.2613
 Epoch 12/50
 191/191 3s 15ms/step -
 accuracy: 0.9828 - loss: 0.0366 - val_accuracy: 0.8004 - val_loss: 2.2819
 Epoch 13/50
 191/191 3s 15ms/step -
 accuracy: 0.9858 - loss: 0.0339 - val_accuracy: 0.8004 - val_loss: 2.2583
 Epoch 14/50
 191/191 3s 15ms/step -
 accuracy: 0.9817 - loss: 0.0353 - val_accuracy: 0.7965 - val_loss: 2.1566
 Epoch 15/50
 191/191 3s 15ms/step -
 accuracy: 0.9839 - loss: 0.0350 - val_accuracy: 0.8004 - val_loss: 2.2722
 Epoch 16/50
 191/191 3s 15ms/step -
 accuracy: 0.9850 - loss: 0.0309 - val_accuracy: 0.7919 - val_loss: 2.3638
 Epoch 17/50

191/191 3s 15ms/step -
 accuracy: 0.9874 - loss: 0.0307 - val_accuracy: 0.7965 - val_loss: 2.2539
 Epoch 18/50
 191/191 3s 15ms/step -
 accuracy: 0.9843 - loss: 0.0321 - val_accuracy: 0.7991 - val_loss: 2.1848
 Epoch 19/50
 191/191 3s 15ms/step -
 accuracy: 0.9869 - loss: 0.0292 - val_accuracy: 0.7919 - val_loss: 2.0238
 Epoch 20/50
 191/191 3s 15ms/step -
 accuracy: 0.9862 - loss: 0.0365 - val_accuracy: 0.7997 - val_loss: 2.1166
 Epoch 21/50
 191/191 3s 15ms/step -
 accuracy: 0.9828 - loss: 0.0323 - val_accuracy: 0.8024 - val_loss: 2.0643
 Epoch 22/50
 191/191 3s 15ms/step -
 accuracy: 0.9852 - loss: 0.0320 - val_accuracy: 0.7945 - val_loss: 2.1305
 Epoch 23/50
 191/191 3s 15ms/step -
 accuracy: 0.9855 - loss: 0.0346 - val_accuracy: 0.8011 - val_loss: 2.2196
 Epoch 24/50
 191/191 3s 15ms/step -
 accuracy: 0.9853 - loss: 0.0305 - val_accuracy: 0.8004 - val_loss: 2.2288
 Epoch 25/50
 191/191 3s 15ms/step -
 accuracy: 0.9809 - loss: 0.0421 - val_accuracy: 0.7971 - val_loss: 2.1823
 Epoch 26/50
 191/191 3s 15ms/step -
 accuracy: 0.9873 - loss: 0.0291 - val_accuracy: 0.8030 - val_loss: 2.2199
 Epoch 27/50
 191/191 3s 15ms/step -
 accuracy: 0.9859 - loss: 0.0310 - val_accuracy: 0.8063 - val_loss: 2.2261
 Epoch 28/50
 191/191 3s 15ms/step -
 accuracy: 0.9856 - loss: 0.0277 - val_accuracy: 0.8017 - val_loss: 2.2722
 Epoch 29/50
 191/191 3s 15ms/step -
 accuracy: 0.9844 - loss: 0.0294 - val_accuracy: 0.8024 - val_loss: 2.3347
 Epoch 30/50
 191/191 3s 15ms/step -
 accuracy: 0.9838 - loss: 0.0353 - val_accuracy: 0.8070 - val_loss: 2.1659
 Epoch 31/50
 191/191 3s 15ms/step -
 accuracy: 0.9866 - loss: 0.0295 - val_accuracy: 0.8076 - val_loss: 2.2348
 Epoch 32/50
 191/191 3s 15ms/step -
 accuracy: 0.9863 - loss: 0.0294 - val_accuracy: 0.8096 - val_loss: 2.1884
 Epoch 33/50

191/191 3s 15ms/step -
 accuracy: 0.9847 - loss: 0.0322 - val_accuracy: 0.8089 - val_loss: 2.2392
 Epoch 34/50
 191/191 3s 15ms/step -
 accuracy: 0.9846 - loss: 0.0307 - val_accuracy: 0.8102 - val_loss: 2.2712
 Epoch 35/50
 191/191 3s 15ms/step -
 accuracy: 0.9862 - loss: 0.0271 - val_accuracy: 0.8089 - val_loss: 2.1921
 Epoch 36/50
 191/191 3s 15ms/step -
 accuracy: 0.9813 - loss: 0.0351 - val_accuracy: 0.8050 - val_loss: 2.1654
 Epoch 37/50
 191/191 3s 15ms/step -
 accuracy: 0.9862 - loss: 0.0312 - val_accuracy: 0.8070 - val_loss: 2.1702
 Epoch 38/50
 191/191 3s 15ms/step -
 accuracy: 0.9837 - loss: 0.0312 - val_accuracy: 0.8056 - val_loss: 2.1213
 Epoch 39/50
 191/191 3s 15ms/step -
 accuracy: 0.9860 - loss: 0.0293 - val_accuracy: 0.8076 - val_loss: 2.2364
 Epoch 40/50
 191/191 3s 15ms/step -
 accuracy: 0.9842 - loss: 0.0293 - val_accuracy: 0.8050 - val_loss: 2.1667
 Epoch 41/50
 191/191 3s 15ms/step -
 accuracy: 0.9847 - loss: 0.0329 - val_accuracy: 0.8043 - val_loss: 2.2310
 Epoch 42/50
 191/191 3s 15ms/step -
 accuracy: 0.9867 - loss: 0.0296 - val_accuracy: 0.8030 - val_loss: 2.3100
 Epoch 43/50
 191/191 3s 15ms/step -
 accuracy: 0.9845 - loss: 0.0443 - val_accuracy: 0.8076 - val_loss: 2.1636
 Epoch 44/50
 191/191 3s 15ms/step -
 accuracy: 0.9852 - loss: 0.0308 - val_accuracy: 0.8063 - val_loss: 2.1197
 Epoch 45/50
 191/191 3s 15ms/step -
 accuracy: 0.9846 - loss: 0.0313 - val_accuracy: 0.8037 - val_loss: 2.1498
 Epoch 46/50
 191/191 3s 15ms/step -
 accuracy: 0.9836 - loss: 0.0356 - val_accuracy: 0.8050 - val_loss: 2.1191
 Epoch 47/50
 191/191 3s 15ms/step -
 accuracy: 0.9852 - loss: 0.0348 - val_accuracy: 0.8056 - val_loss: 2.1753
 Epoch 48/50
 191/191 3s 15ms/step -
 accuracy: 0.9826 - loss: 0.0311 - val_accuracy: 0.8050 - val_loss: 2.2017
 Epoch 49/50

191/191 3s 15ms/step -
accuracy: 0.9853 - loss: 0.0304 - val_accuracy: 0.8043 - val_loss: 2.1776
Epoch 50/50
191/191 3s 15ms/step -
accuracy: 0.9841 - loss: 0.0295 - val_accuracy: 0.8037 - val_loss: 2.1894



48/48 0s 5ms/step
F1 Score Combined Increased Epoch Size/Decrease Learning Rate:
0.7722772277227723

Hyperparameter Tuning Results

The table below shows the performance metrics for the different hyperparameter tuning strategies. These hyperparameter modifications did not appear to have a significant outcome on the accuracy or f1 score of the validation data. The accuracy of the training set is very good and the validation accuracy and f1 of 0.8 and 0.77 respectively are not too bad, but the differences suggest there is still inherent overfitting in the model.

Accuracy		
	Train	Val
Initial Model	0.9839	0.8017
Decrease Learning Rate	0.9765	0.7945

Accuracy		
Switch Optimizer	0.9834	0.7991
Increase Batch Size	0.9821	0.7971
Increase Epochs	0.9864	0.7979
Combined	0.9841	0.8037

F1 Score	
	Val
Initial Model	0.7705
Decrease Learning Rate	0.7659
Switch Optimizer	0.7650
Increase Batch Size	0.7610
Increase Epochs	0.7656
Combined	0.7723

1.1.6 Conclusion

This project provided a good exercise in NLP (Natural Language Processing) and also showed the potential in creating RNN/LSTM models. The model I created was able to classify whether a twitter tweet was about an actual disaster with approximately 80 percent accuracy. In this project I experimented with tuning hyperparameters to get an even better accuracy, however these results did not lead to too much improvement. For future modifications in attempting to achieve higher accuracy, making changes to the model's architecture such as adding additional layers may be needed. Making modifications to the text processing could also be an improvement to explore in the future.

1.1.7 References

AnmolsX_test_seq0, February 7). Disaster Tweets : Simple RNN Implementation. Kaggle.com; Kaggle. <https://www.kaggle.com/code/anmolstha/disaster-tweets-simple-rnn-implementation>

Andreshg. (2021, May 10). NLP GloVe, BERT, TF-IDF, LSTM... Explained. Kaggle.com; Kaggle. <https://www.kaggle.com/code/andreshg/nlp-glove-bert-tf-idf-lstm-explained>

Feldges, C. (2022, April 2). Text Classification with TF-IDF, LSTM, BERT: a quantitative comparison. Medium. <https://medium.com/@claude.feldges/text-classification-with-tf-idf-lstm-bert-a-quantitative-comparison-b8409b556cb3>

Kumar, V. (2020, February 23). Real or Not? NLP with Disaster Tweets (A Data science Capstone Project). Medium; Real or Not? NLP with Disaster Tweets. <https://medium.com/real-or-not-nlp-with-disaster-tweets/real-or-not-nlp-with-disaster-tweets-a-data-science-capstone-project-fafa6c35c16f>

```
[316]: ###
# Kaggle Submission Code - score 0.779
###
y_pred = model.predict(x_sub_test)
y_pred = np.round(y_pred).astype(int).reshape(3263)
```

```

print(y_pred)

sub = pd.DataFrame(
    list(zip(df_test['id'], y_pred)),
    columns=["id", "target"],
)
print(sub)
sub.to_csv("submission.csv", index=False)

```

102/102 0s 4ms/step

[0 1 0 ... 1 1 0]

	id	target
--	----	--------

0	0	0
---	---	---

1	2	1
---	---	---

2	3	0
---	---	---

3	9	1
---	---	---

4	11	1
---	----	---

...
-----	-----	-----

3258	10861	1
------	-------	---

3259	10865	1
------	-------	---

3260	10868	1
------	-------	---

3261	10874	1
------	-------	---

3262	10875	0
------	-------	---

[3263 rows x 2 columns]

[]: